

10/10 answer to an AI Product Sense question 2

Nail your AI Product Sense interview round



RIDHIMA KHURANA

MAR 10, 2026 · PAID




9



1


Share



This was recently asked at a number of AI-native companies, including OpenAI. You can apply this framework to numerous other questions around AI impacting productivity, buying art, experiencing travel, etc.

Let's get into it!

“Imagine the year is 2035. Design a new consumer product that leverages advances in AI to fundamentally change how people learn.”



1. Comprehend: Goals and Assumptions

Alright, let's get into this. I love this question because the word *fundamentally* is doing an enormous amount of heavy lifting. The interviewer isn't asking you to build Duolingo with better animations. They're asking you to question the act of learning itself—which has been running on the same basic operating system since the 1800s: one teacher, thirty students, same pace, same test. We digitized the textbook. We never reimagined the classroom.

Coaching note: Notice how I'm breaking the prompt apart before jumping into solutions. When an interviewer uses a word like “fundamentally,” they're testing whether you can hear what the question is *really* asking. Also—what does “learn” mean here? Anything. Academics, skills, hobbies, instruments. Acknowledge that breadth and then narrow it later.

Here's how you'd walk through it:

“First, let me anchor this in a clear **Mission**. If I'm at a company like OpenAI or Anthropic in 2035, we aren't chasing market share in EdTech—we're trying to make

mastery of any skill accessible to any human, regardless of starting point, income, or geography. We're democratizing *expertise*, not distributing content. If we get this wrong, we risk widening the gap between people who can afford great learning and people who can't. So Equity is our 'Feature Zero.'"

Why this works: You've immediately told the interviewer you think about second-order consequences. That's the kind of thinking AI-native companies are screening for.

"Now, for some **Technical Context**. I want to ground this in what AI can realistically do in 2035, because if I hand-wave the tech, this answer falls apart. Here are my assumptions:

Multimodal Reasoning has matured. Models can watch you *attempt* a task—a piano chord, a negotiation, a calculus proof—and diagnose *why* you're stuck, not just *that* you're stuck.

Long-Context Personal Memory means an AI holds years of your learning history in working context. Think of it as a tutor who has been with you since kindergarten and never forgot a single lesson.

Real-Time Cognitive Load Sensing through lightweight wearables—eye tracking, micro-expression analysis, galvanic skin response. The AI can read when you're in flow, when you're frustrated, and when you're three seconds from closing the laptop and opening Netflix.

Generative Simulation has matured. You don't *watch a video* about giving difficult feedback to a direct report; you *give the feedback* to a simulated version of that person who adapts in real time."

Coaching note: Laying out your tech assumptions does two things: (1) shows the interviewer you actually understand AI capabilities, and (2) makes every solution you propose later feel *earned* rather than magical.

“Now let’s talk about the **Weaknesses**. AI still struggles with teaching *wisdom*—the kind of judgment that comes from twenty years of lived experience. And there’s the ‘Crutch Effect’: if the AI is too helpful, learners never develop what researchers call ‘productive struggle’—that uncomfortable zone where real learning actually happens. We have to design *for* that tension, not around it.”

See what happened there? You just named a constraint that will show up later as a product principle. Plant seeds early that bloom into design decisions later.

“Looking at **Macro Trends**:

The Skills Half-Life Crisis: By 2035, a professional skill lasts roughly 2-3 years before obsolescence. Continuous reskilling is oxygen.

The MOOC Graveyard: Online courses proved content alone isn’t enough—completion rates hover around 5-15%. The world doesn’t have an information shortage; it has a *relationship* shortage. Content without connection is just a library nobody visits.

The Credential Collapse: Employers care about *what you can do*, not *where you studied*. Skills-based hiring creates a massive opening for a product that *verifiably demonstrates* competence.


AR/VR at Consumer Price Points: Lightweight AR glasses have gone mainstream. Immersive practice environments are an everyday tool.”

Coaching note: Each trend directly supports a design decision you’ll make later. Skills Half-Life Crisis justifies your *market*. MOOC Graveyard justifies your *approach*. Credential Collapse justifies your *value proposition*. AR/VR justifies your *form factor*. None of these are decorative—they’re structural.

“Finally, the elephant in the room: **The ‘Cognitive Surveillance’ narrative**. Any AI that deeply understands how your brain learns—your weaknesses, your emotional triggers

—holds extraordinarily sensitive data. A data breach here doesn’t just leak your password; it leaks your incompetencies. To mitigate this, we need gold-standard ‘Cognitive Privacy’ protocols—on-device-only processing as the default—before we ship a single unit.”

Now that we’ve set the stage, let’s talk about who this is actually for.



2. Identify Users: Who Are the Key Users/Personas? How Would You Prioritize Your Top Persona and Why?

“When you have a technology this powerful, the temptation is to say ‘everyone,’ but in a GTM with ethical stakes this high, we need to be surgical. I’ve mapped out three primary ecosystems:”

Coaching note: Present three segments, explain why each is compelling, then explicitly kill two with conviction. The interviewer is testing whether you can *kill your darlings*.

The Three Segments

- 1. K-12 / Formal Education** — The largest market. Every parent fantasizes about a private tutor. This technology is that tutor—endlessly patient, perfectly adaptive, available at 2 AM.
- 2. Enterprise L&D** — \$380 billion annually on training, most of it PowerPoint deck employees click through while checking their phones. The ROI story writes itself.
- 3. The Ambitious Self-Reinventor** — A 28-42-year-old professional with a specific, high-stakes skill gap causing *active pain*. A brilliant engineer who just got promoted to management and can't give feedback without starting a small emotional fire. A founder whose Series A pitch currently puts investors to sleep. They aren't learning fun. Something in their life *depends* on this.

Which Segment and Why


Deprioritizing K-12: Procurement bureaucracies that move at the speed of a glacier with a limp. COPPA regulations. Teacher union dynamics. And if our model has a 'Learning Hallucination'—confidently teaching a child something subtly wrong—the consequences for a developing mind are severe. It's a North Star, but not our MVP.

Deprioritizing Enterprise: This is the 'Cognitive Surveillance' danger zone. The moment a company asks an employee to learn with an AI that tracks cognitive weaknesses, we've created a dystopian performance review nightmare. One leaked 'Learning Profile' and our brand is finished.

Targeting The Self-Reinventor: *Intrinsic motivation*—they chose this, so they'll engage deeply, giving our AI the rich signal it needs to improve. *Quantified Self mindset*—they already spend \$50-300/month on self-improvement. *AI tolerance*—they'll treat a weak recommendation as a data point, not a verdict. *High feedback signal*—gold for model fine-tuning.

I see this as a **D2C Subscription**. We aren't selling an AI tutor; we're selling a private ecosystem for skill transformation. Start with this motivated, tech-literate group, refine accuracy, then scale to more vulnerable populations.

Coaching note: See how each deprioritized segment ties back to a risk from Section K-12 connects to "Learning Hallucination." Enterprise connects to "Cognitive Surveillance." That's how you make an answer feel like one interconnected argument



3. Needs, Challenges, and Product Principles

“So, we’ve pinpointed our user. What is the actual job-to-be-done?”

The core need is what I’d call the ‘**Transfer Bridge**.’ They feel a persistent gap between what they *know* and what they can *do* under pressure. They’ve watched the YouTube videos. They can explain the theory of giving great feedback. And then their direct report gets defensive and they freeze up and say ‘yeah, no, everything’s fine actually.’ The gap between *knowing* and *doing* is where 90% of learning investments go to die.’

See what I did there? I coined a term—'Transfer Bridge.' It gives the interviewer a mental anchor. Every time you reference it later, they remember the vivid image. Coined terms are your best friend in PM interviews.

"But to build that bridge, we have to solve the **Crux**: Practice without context is theater. This is where we hit the 'Black Mirror' challenges:

The 'Learning Hallucination': The AI confidently coaches you on a communication approach that's subtly wrong for your cultural context. A *social* error coached by AI could damage a real relationship.

The 'Crutch Spiral': The AI intervenes too early, and you never develop resilience. It could accidentally create the most well-informed generation of helpless professionals the world has ever seen.

Cognitive Data Sovereignty: Your learning profile—every gap, every struggle—is profoundly intimate data. A hack doesn't leak your credit card; it leaks a map of everything you're bad at.


Now—the satisfying part—each risk becomes a product principle:

Cognitive Sovereignty. Privacy isn't a feature; it's the product. On-device-only processing. If the data doesn't exist in the cloud, it can't be subpoenaed, hacked, or sold.

Struggle by Design. The AI keeps you in the zone of 'desirable difficulty'—slightly beyond your ability, where neural rewiring happens. It steps in only when you tip from 'challenged' into 'demoralized.' Think of it as the world's most perceptive climbing belayer: it lets you struggle, but catches you before you hit the ground.

Hypotheses over Verdicts. The AI never says, 'You are bad at leadership.' It says, 'I noticed your confidence drops when conversations turn confrontational—want to explore that?' We're offering mirrors, not report cards."

Coaching note: The rhythm is Risk → Principle → Vivid metaphor. “Crutch Spiral” “Struggle by Design” → “climbing belayer.” This makes principles feel inevitable rather than arbitrary.



4. Solutions

“Let’s get into the actual product. Three distinct ways this could manifest:

1. AI-Generated Practice Simulations

You tell the AI: ‘I need to tell my co-founder their architecture is creating tech debt and they get defensive.’ The AI generates a photorealistic, interactive simulation—your co-founder’s avatar in AR, reacting dynamically. After each round, it debriefs: 2:14 when they pushed back, your language shifted from assertive to apologetic. What to try a different approach at that exact moment?’

The Goal: Attack the Transfer Bridge directly. You aren’t learning *about* feedback; you’re *giving* feedback. You can fail spectacularly and nobody gets hurt.

2. Adaptive Learning Sequencing

Instead of a fixed curriculum, the AI builds a dynamic path unique to you. It maps your competencies, identifies gaps, and sequences micro-sessions (10-15 min) matching your cognitive style. If you're an engineer learning leadership, it explains stakeholder management through systems architecture metaphors.

The Goal: Solve the Pacing Problem. Every session feels designed for you. Because it was.

3. AI-Facilitated Peer Learning Cohorts

The AI matches you with 3-4 learners at a similar stage. It facilitates weekly peer-coaching sessions, noticing things humans miss: 'Priya, you naturally used Socratic confrontation when giving feedback to Marcus—would you like to make that a deliberate tool?'

The Goal: Solve the Loneliness Tax. MOOCs have a 5% completion rate. CrossFit has 90% retention. The difference isn't content—it's community.


Prioritizing the MVP

Leading with **Solutions 1 + 2**. Why not peer cohorts? Matching requires critical mass—chicken-and-egg problem for Day 1. It's our Phase 2 retention play.

Why simulations first? It's the 'holy sh—' moment. The first time you rehearse a real conversation, fail safely, get coaching, and walk away feeling *ready*—the product justifies its existence in a single session. That's the story users tell friends at dinner.

Critical tradeoff: Depth over breadth. MVP focuses on **Leadership Communication** only—giving feedback, navigating conflict, presenting ideas. It's the biggest skill gap for our persona, inherently practice-based, and where simulation shines brightest."

Coaching note: Every prioritization links back to something earlier. Deprioritizing with *respect*—not dismissal—shows maturity.



5. MVP, Tradeoffs, and Metrics

“Our MVP: a D2C Subscription with AR-enabled simulations and personalized adaptive sequencing. How do we know if we’re winning?”

Primary: Real-World Transfer Rate. ‘Did you attempt this skill in a real situation this week? How did it go?’ If users aren’t transferring to the real world, we’re just a very expensive video game.

Secondary: Time-to-Confidence. How many sessions until a user reports feeling confident? Traditional methods take 3-6 months for soft skills. Our target: a measurable shift within 4 weeks.

Engagement: Weekly Active Practice Sessions. Not ‘time in app’—that rewards doomscrolling. We measure *deliberate practice sessions*. Target: 3+ per week.

Safety: The ‘Distress Signal.’ We track users who exit mid-simulation and don’t return for 72+ hours. If the distress rate exceeds 5%, we recalibrate the simulation’s intensity.

Coaching note: That safety metric separates good from great. Most candidates stop at “I’m watching for how this product could hurt someone.” At AI-native companies, that’s the line that gets you the offer.

The Close

Now, you always want to end with an “Imagine” statement. This is your mic drop. It should be visceral, human, and emotionally resonant. You want the interviewer to feel the product, not just understand it:

“Imagine a world where a first-time manager doesn’t have to learn leadership by accidentally damaging three relationships and losing one great employee. Where a founder doesn’t have to bomb four fundraising pitches before figuring out how to tell their story. Where a brilliant introvert can rehearse the hardest conversation of their career—ten times, in private, with an AI that knows their exact patterns—and walk into the real room feeling, for the first time, like the person they’ve always wanted to be.

We aren’t just building a learning product. We’re closing the oldest gap in human education: the gap between *knowing* and *doing*.”

COACHING NOTE

Frameworks get you in the door.
First-principles thinking
gets you the offer.

I offer 1:1 coaching for AI-first PM interviews — OpenAI, Anthropic, DeepMind, and the next frontier lab.

If you want to learn how to think this deeply and handle any question thrown at you in AI-PM interviews—whether it’s OpenAI, Anthropic, DeepMind, or the next frontier lab—I offer 1:1 coaching to help you build this muscle. Frameworks get you in the door; first-principles thinking gets you the offer.



9 Likes • 1 Restack

Discussion about this post

Comments

Restacks



Write a comment...

© 2026 Ridhima Khurana · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture