10/10 Answer to "Paint a Vision for a New Anthropic Product"

How to tackle the most open-ended question in an AI PM interview — and why mos candidates drown in possibility.

RIDHIMA KHURANA
APR 01, 2026 · PAID

6

Share

Every product sense question has a trap. "Improve Instagram Stories" traps you int incremental thinking. "Design a product for Mars colonists" traps you into sci-fi fantasy. But "Paint a vision for a new Anthropic product" has a trap that's sneakier than both of those combined.

It traps you into freedom.

When everything is on the table, most candidates do one of two things. They either panic and pitch something painfully safe ("Claude, but for email!") or they go full El and pitch something so wildly ambitious that the interviewer mentally checks out b sentence three. The sweet spot is a product that feels inevitable — something wher the company's unique DNA, the market timing, and an underserved user base all sn together like a seatbelt clicking into place.

That's what we're going to build today.

I'm going to walk you through exactly how you'd answer this in an Anthropic interview, using the product sense framework. We're going to think deeply, coin a f terms, and by the end, you'll have a product vision that could make Dario Amodei pause his morning cortado.

Let's get into it.

1. Comprehend: Goals and Assumptions

“Alright, this is a wide-open prompt, so the first thing I want to do is narrow the aperture. And the way I do that is by asking: what is Anthropic *uniquely* positioned build that nobody else can?

Let me start with the **Mission**. Anthropic exists to build AI that is safe and benefici for humanity. That phrase — ‘safe and beneficial’ — is doing a lot of heavy lifting h It means the *right* product idea for Anthropic has safety baked into its DNA, where their alignment research becomes a product advantage rather than a constraint. I’n looking for a domain where trust is oxygen.

Now, **Strengths**. Anthropic’s moat is Constitutional AI and the safety-first brand they’ve built. Claude is widely considered the most ‘trustworthy’ frontier model in t market. Cowork is expanding their agent footprint. They’ve got deep research talen and recently launched Claude Code Security, which shows they can turn safety expertise into a shipped product. They also just committed to keeping Claude ad-fr — which signals something important about how they think about user relationshi

Weaknesses to keep honest about: Anthropic has a smaller consumer distribution footprint than Google or OpenAI. No consumer hardware. No built-in ecosystem li Apple or Google has with schools. So whatever we build, the GTM needs to work without a pre-existing distribution moat.

Let me look at the **Macro Trends**, because this is where the product starts to reveal itself:

The ‘Unsupervised AI’ Crisis with Kids. Here’s the elephant stampeding through t room: kids are already using AI. Heavily. Studies show 86% of students use AI in th studies. A lot of kids are using chatbots to deal with anxiety and personal issues. Bu the thing is that they’re using tools that were built for adults, with no guardrails designed for developing minds. It’s like giving a twelve-year-old the keys to a Form One car and saying ‘have fun, the brakes are somewhere.’

The Teacher Shortage. There are over 300,000 unfilled teaching positions in the US alone. The ones who are there are drowning — trying to differentiate instruction across 30 kids with wildly different levels. They don’t need another LMS. They need force multiplier.

The ‘Homework Machine’ Backlash. Schools are in a tug-of-war with AI. Some are banning it. Others are embracing it. But almost nobody has cracked the code on AI that actually *teaches* rather than just *tells*. The entire industry is solving the wrong problem: they’re optimizing for ‘getting the answer’ when the value is in ‘getting th struggle.’

The Parental Anxiety Spiral. Parents know their kids are using AI. They can’t stop And they feel completely out of the loop. This is a massive, unaddressed emotional need.

So here's where I land: **Anthropic should build Claude Learning Companion — an tutor designed specifically for kids aged 10–14, where the core product promise is that it teaches thinking, never gives answers, and keeps parents in the loop.**

Why this product, for this company, right now? Because Anthropic is the only front AI lab with the safety credibility to walk into this space and be *believed*. When Goog says 'safe AI for kids,' parents hear 'the company that monetizes my data wants acc to my child's brain.' When Anthropic says it, it lands differently — because the enti brand has been built on the promise that safety comes first.

And **Feature Zero is 'Never Replace the Adult in the Room**.' The AI doesn't becom the teacher or the parent. It becomes the incredibly patient study partner sitting ne to the kid at the kitchen table that asks really good questions and never, ever does t homework for them."

> **Coaching note:** See what happened there? I didn't just list trends — I *connected* them. Each macro trend directly supports a design decision I'll make later. The teacher shortage justifies why AI help is needed. The 'Homework Machine' backlash tells me what the product must *avoid*. The parental anxiety tells me wh the *real* buyer is. In an AI-native interview, your Comprehend section should rea like a thesis, where each observation earns its seat at the table. Also notice: I coi 'Feature Zero' early. That's a memory hook. The interviewer will use that phrase their debrief. Make their job easy.

2. Identify Users: Who Is This Actually For?

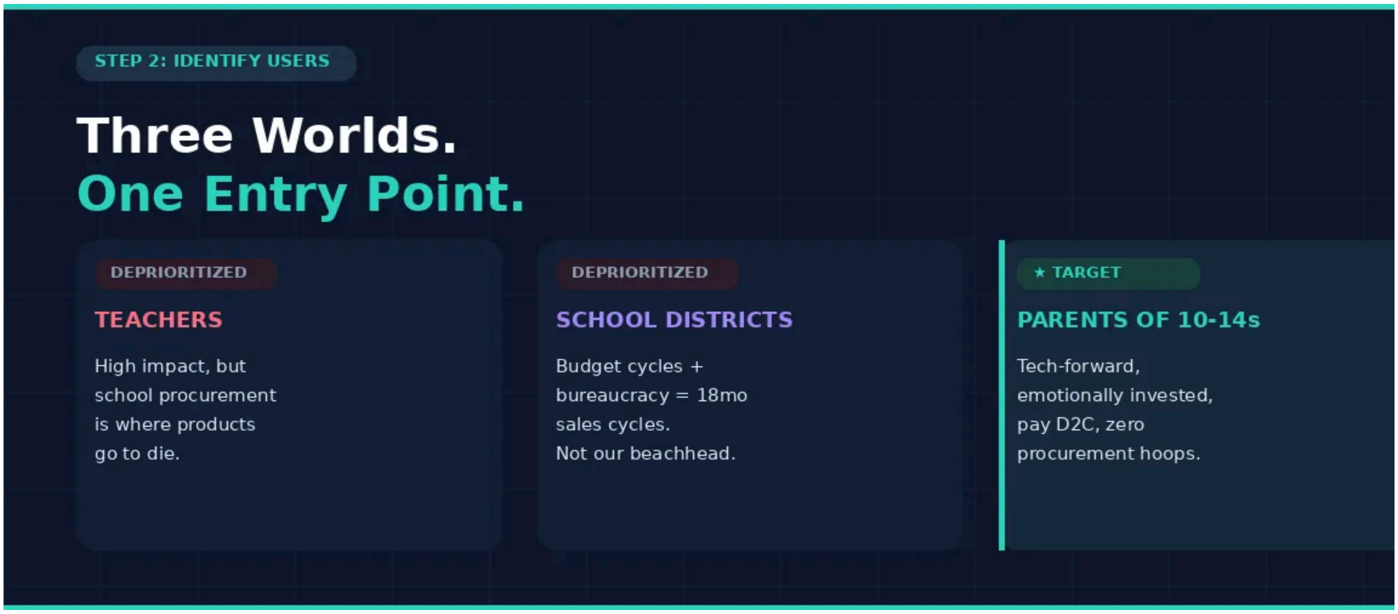

"Now that we've established the 'why,' let's talk about the 'who.' With a product this sensitive — we're talking about children — the temptation is to design for the obvi user (the kid) and ignore the ecosystem. That would be a fatal mistake. I see three distinct user worlds here:

1. Teachers

Why they matter: Teachers are the credibility layer. If teachers trust the product, it gets into classrooms, and that's the ultimate scale play. A great teacher integration could mean this shows up in lesson plans, gets recommended at parent-teacher conferences, becomes word-of-mouth gold.

Why I'm deprioritizing them for launch: School procurement is where good produc go to die. You're talking about district-level budget approvals, a patchwork of state-level data privacy laws, 12–18 month sales cycles, and archaic IT departments. We'd burn through our runway getting one pilot while a competitor ships to a million homes. Teachers are our expansion play, not our MVP audience.

2. School Districts / Administrators

Why they matter: If we could land a district partnership, we'd have instant distribut to thousands of students and a recurring revenue engine.

Why I'm deprioritizing: Same procurement problem as teachers, but multiplied. District buyers optimize for cost and compliance, not pedagogical innovation. They want complaice, SOC 2 reports, integrations with Google Classroom and Canvas. A necessary but also slow. And the political risk is enormous: one superintendent championing 'AI for kids' becomes a campaign attack ad during the next school bo election. The channel is real but the timing is premature.

3. Parents of 10–14 Year Olds (The Target)

This is our precision strike. Let me paint the persona. They're parents of a 12-year- who just started pre-algebra and a sinking feeling that ChatGPT is doing half the homework. They're tech-forward and probably use Claude themselves for work. Th *understand* AI. They know they can't ban it. But they want to *channel* it.

Why this specific age range? Because 10–14 is the crucible. It's when abstract reasoning kicks in, when the curriculum gets genuinely hard, and when kids start doing homework independently (unsupervised with a laptop). It's also the age wher learning habits calcify. Get a kid thinking well at 12, and you've shaped how they approach problems for life.

Why D2C through parents? Three reasons.

1. **Zero procurement friction**: Parents have a credit card and can sign up tonight.
2. **Emotional investment**: They're not buying a SaaS tool, they're buying peace of mind about their child's intellectual development. That's a fundamentally diffe willingness-to-pay. T
3. T**he trust bridge**: By making parents the buyer *and* the oversight layer, we solve the 'who's watching the AI' problem at the product level rather than the policy level."

> **Coaching note:** Notice I deprioritized two perfectly valid user segments *and explained why with specificity*. This is where most candidates fall apart on open-ended questions. They pick a user and move on. Strong candidates show the interviewer the segments they *rejected* and the reasoning behind each cut. It demonstrates you're not just picking randomly, you're making a strategic bet.

3. Needs, Challenges, and Product Principles

"So, we've locked in on Parent and their 12-year-old. Now let's get under the hood. What do parents actually *need*, and what are the land mines between us and a produ they'd trust with their kid?

The Core Job-to-Be-Done is what I'd call **The Guided Struggle.** Parent don't want that makes their kid's life easier. They want AI that makes their kid's *thinking* sharp They wants the educational equivalent of a climbing wall with a really good spotter who lets the kid stretch, sweat, even slip a little, but never lets them fall.

That sounds simple. It is brutally hard to build. Here are the challenges:

Challenge 1: The 'Just Tell Me' Problem. Kids are smart optimizers. The first thing 12-year-old will try is 'just give me the answer.' Every existing AI does exactly that. our product caves to that pressure even once, the trust contract with the parents is broken. We need the AI to be constitutionally incapable of giving direct answers to homework questions. This isn't a prompt engineering hack. This needs to be an architectural commitment, built on Anthropic's Constitutional AI approach.

Challenge 2: The Calibration Tightrope. If the AI helps too much, the kid learns nothing. If it helps too little, the kid gets frustrated, rage-quits, and tells parents th app is 'stupid.' The product has to dynamically calibrate how much scaffolding to o based on the kid's demonstrated understanding and it needs to get this right within

the first 3–4 interactions. We need to get to the point where productive difficulty ti
into destructive frustration. Getting this wrong in either direction is a product kille

Challenge 3: The Transparency Deficit. Parents right now have zero visibility into what their kids do with AI. It’s a complete black box. They might see the finished homework and think their kid is thriving, when in reality, Claude (or ChatGPT) wrc the whole thing. Our product has to solve this at the *structural* level with a proactive communication layer that gives parents real-time insight into what the kid is struggling with, where they’re growing, and where the AI had to step in.

To navigate these challenges, here are the **Product Principles:**

Principle 1: The Socratic Guardrail. The AI never gives answers. It asks questions, offers hints, provides analogies, and celebrates breakthroughs, but the final cogniti leap always belongs to the kid. This is our ‘Feature Zero.’ We lean into Anthropic’s Constitutional AI strength to make this architecturally, not just behaviorally, enfor

Principle 2: Productive Friction by Design. We resist the instinct to make everythi frictionless. The struggle IS the product. We design for the number of genuine moments of understanding per session. If a kid breezes through without effort, we’ failing.

Principle 3: The Parent is the Product. Parents are not a secondary stakeholder wh get a monthly email report. They’re a co-user. The transparency layer is designed sc well that it *itself* becomes a reason to subscribe because it gives them something they’ve never had: a clear, honest view of how their child thinks.

> **Coaching note:** I just introduced a coined term: ‘Socratic Guardrail” This is intentional. Coined terms do two things in an interview: they signal that you’re thinking from first principles (you’re naming concepts, not borrowing them), anc they give the interviewer a vocabulary to retell your answer. When the debrief happens, your interviewer won’t say ‘they talked about scaffolding difficulty.’ Th

> say ‘they had this concept called the “Socratic guardrail” That’s stickiness. Also notice how each challenge maps directly to a product principle. That’s not coincidence — it’s structure. The challenges create tension; the principles resolv it.

4. Solutions

“Now, let’s translate those principles into actual products. I’ve brainstormed three distinct capabilities:

1. The Question Engine (Socratic AI Core)

When a kid is stuck on a pre-algebra problem, he /she doesn’t get an answer. They g a question. And then another question. And then, when they’re close, maybe an analogy.

For example, the kid types: ‘What is 3x + 7 = 22?’ The Question Engine responds: ‘Interesting. If I told you that you have 22 apples, and 7 of them are in a basket, how many are left on the table? Now… if those remaining apples were split into 3 equal groups…'

AI needs to model the *student's* current understanding in real time. It's using Anthropic's reasoning capabilities to figure out not just 'what is the right answer' b 'what does *this specific kid* misunderstand right now, and what's the smallest nudge would unblock them?' That's a genuinely hard AI problem and it's one where Anthropic's model capabilities give us a real edge.

2. The Struggle Meter (**Adaptive Difficulty Calibration**)

This is the product's nervous system. It tracks three signals across every session: tir on-problem, error patterns, and emotional cues from the language the kid uses (thir like 'this is stupid' vs. 'wait, let me try again').

When the kid is in the productive struggle zone, making mistakes but trending tow understanding, the Meter stays quiet. When they cross the Struggle Threshold into genuine frustration, it adjusts: maybe it offers a worked example of a *similar* (but different) problem, or breaks the current problem into sub-steps. And crucially, it *n* adjusts down preemptively. The default assumption is that the kid can handle it.

3. The Parent Window (**Transparency Dashboard**)

This is the parents product. After each study session, they get a 3-sentence summar that tells them: what their kid worked on, what he/she struggled with, and where he/she had a breakthrough. Not a grade. Not a score. A *narrative*.

Prioritizing the MVP:

I'm leading with the **Question Engine** as the MVP because:

The Struggle Meter is essential but technically complex, getting the calibration rigl requires significant training data from actual student interactions, which we don't have yet. We need to ship and learn first.

The Parent Window is a retention and trust play, but it’s meaningless without a cor learning experience that’s actually good. You don’t need a dashboard for a product nobody uses.

The Question Engine is the atomic unit. It’s where the core value proposition lives- that teaches your kid to think.’ If this works, everything else accelerates. If it doesn’ nothing else matters.

For the MVP, I’d scope it to **math for ages 10–14**. Because math has objectively rigl answers (which makes the Socratic approach testable), a clear progression of conce (which makes the scaffolding solvable), and it’s the subject parents worry about mo It’s also the subject where ‘just tell me the answer’ is most destructive — because m builds on itself. If you don’t understand fractions, algebra will eat you alive.

We package this as a D2C subscription: $19.99/month. The parent signs up, creates supervised account for their kid, and gets a weekly Learning Pulse from day one.

> **Coaching note:** See how I brainstormed three solutions, then cut to one for the MVP? This is the ‘show your range, then show your judgment’ move. The interviewer wants to see that you can ideate broadly AND prioritize ruthlessly. I also scoped the MVP to one subject (math) and one age range (10–14). Specificity confidence. When a candidate says ‘we’d start with all subjects for all kids,’ the interviewer hears ‘I don’t know how to make hard calls.’ When you say ‘math, ag 10–14, $19.99/month,’ they hear someone who could actually ship this.

5. MVP, Tradeoffs, and Metrics

“Let me tie this together with how we’d measure success. For a product where the c promise is ‘teaches thinking,’ our metrics have to actually capture *thinking.*

Primary Metric: Independent Solve Rate.

After the AI helps the kid understand a concept through Socratic questioning, can l solve the *next* problem of that type on his own? We measure this by presenting a follow-up problem 24 hours later without AI scaffolding.

This is our North Star because it directly measures the promise we’re making to parents: ‘your kid will actually learn to think.’

Secondary Metric: Parent Trust Score.

Every two weeks, we ask parents a single question: ‘Do you believe this AI is genuir teaching your child, or just helping them get through homework?’ We want this on 1–5 scale, and we want to see it trend upward over the first 90 days. If Trust Score is flat or declining, it means our Learning Pulse and Thinking Map aren’t doing their

But this is a self-reported metric. What could go wrong? Well, for this product, perception IS reality. Even if our independent solve rates are off the charts, if paren *feel* like the AI might be doing the work, they'l cancel. Trust is measured in feelings not dashboards.

Guardrail Metric: Answer Leak Rate.

This is our 'oh no' metric. How often does the AI inadvertently give away an answe despite the Socratic Guardrail? We measure this through a team that runs adversar student prompts ('just tell me,' 'my mom said you can give me the answer,' 'I'll fail i you don't help') and checks whether the AI holds the line. Your classic eval method

This is a product where a leak isn't just a bug, it's a brand-destroying event. One vi TikTok of a kid saying 'I just told Claude my dog died and it gave me all the answer and we're done.

The Tradeoff I'm Making:

I'm deliberately *not* optimizing for session time or daily active usage. That's the old playbook, the engagement trap. If the kid uses the product for 15 focused minutes a solves two problems with genuine understanding, that's better than 45 minutes of A assisted coasting. For a product designed for children, this is both the ethical choic and the strategic one because parents will never trust a product that seems designe keep their kid glued to a screen.

Imagine a world where a 12-year-old sits down with a math problem that makes no sense, a instead of Googling the answer or asking ChatGPT to do it for them, they open an app tha asks them a really good question. And then another. And then — click — they get it. On th own. And across the house, their mom gets a three-sentence note on her phone that says: 'A cracked two-variable equations today. He struggled with isolating y, but figured out he was confusing the subtraction step. He'll see a follow-up tomorrow.'

We're building the first AI product that makes a parent more confident in their child's thinking, and makes a child more confident in their own mind.

That's not just a product. That's a mission statement Anthropic was built to fulfill."

Frameworks get you in the door. First-principles thinking gets you the offer.

If you want to learn how to think deeply and handle *any* question in AI-first intervi
— the open-ended ones, the 'Black Mirror' ones, the ones that don't have a playboo
— I offer 1:1 coaching for PM candidates targeting Anthropic, OpenAI, DeepMind
and other frontier AI companies. [Reach out to learn more.]

6 Likes

Discussion about this post

Comments Restacks

Write a comment...