

The OpenAI PM Interview Process: What Nobody Tells You (Until Now)

And How to Actually Prepare for It



RIDHIMA KHURANA

FEB 24, 2026 • PAID




18



1

Share



I've coached 300+ candidates into FAANG and AI Labs PM roles.


And it's usually the most prepared ones who bomb the OpenAI interviews.

They walk in with their "redesign Google Maps" framework, their CIRCLES method and their beautiful RICE prioritization spreadsheet.

And OpenAI looks at them like... sir, this is a frontier AI lab.

Let's save you from being that person.

First, understand the math.



OpenAI has roughly 30 PMs.

Thirty. Total. Not per team. Not per office. Total.

Meta has over 2,000. let's not try to count the number of Google PM's.

That means every single hire matters. There's no "let's slot them somewhere and see how they do." If you're not a precise fit for the specific role, you're out. You need to pick your role like you're picking a life partner. With intention, research, and maybe slight obsession.


Step 1: Find your role (and actually understand it)

Before you apply, go deep on the product area. Not "I read the blog post" deep. More like "I have opinions about their technical architecture decisions and I've tested 6 competitor tools" deep.

One of my candidates was interviewing for a ChatGPT Enterprise role. Before the recruiter screen, she'd already mapped out the customer segments, tested Copilot for Business and Gemini for Workspace side by side, and had a hot take on why enterprise AI adoption was bottlenecked by IT procurement, not product quality.

She got the offer.

Step 2: The recruiter screen (this is NOT a vibe check)



Most candidates treat this like a scheduling call with some small talk.

Most candidates are wrong.

The recruiter screen at OpenAI is evaluative from minute one. You need to nail three things:

Why OpenAI — specifically. "AI is the future" is not an answer. That's a bumper sticker.

What they want to hear is something like: “I’ve been using the API to build internal tools for my team. I rebuilt our entire customer research workflow using GPT-4, and cut synthesis time by 60%. I went from skeptic to evangelist in three months. I want to be on the team building these capabilities, not just consuming them.”

See the difference? One is a thesis. The other is a TED talk intro.

Your story, mapped to THIS role. Do not do the resume walkthrough. I beg you. “I started at Company A in 2015 where I...” — no. They’ve read your resume. They don’t need the audiobook version.

Instead, pull the single thread that connects your experience to this specific role. If you’re going for a growth PM position, your story should be a growth story. Every bit of it.

Your unique edge. What do you bring that the other 500 applicants don’t? This requires actual self-awareness, which is also something they evaluate in the final round.

Step 3: The product screen

04 — THE PRODUCT SCREEN

FORGET

- "Redesign Google Maps"
- "Improve Instagram Stories"
- "Build a feature for X"

EXPECT

- How would you bring this capability to market?
- What use cases would you prioritize — and deliberately skip?
- What are the risks — including ethical ones?

THE ETHICS PIECE ISN'T A BOX TO CHECK. IT'S A CORE EVALUATION CRITERION.

This is where most traditional PM prep falls apart.

There's no "how would you improve Instagram Stories." There's no existing product redesign.

Instead, you'll get questions like:

- How would you bring [new AI capability] to market?
- What use cases would you prioritize — and what would you deliberately NOT do first?
- What are the risks? And yes, they mean ethical risks too.
- See [my answer to a recent Product Sense question asked at Open AI](#)

The ethics piece isn't a box to check. It's a core evaluation criterion. AI product decisions carry weight that "should we add a dark mode" simply did not.

One candidate I worked with got a question about launching a new multimodal capability. She spent 80% of her answer on the go-to-market and 20% on risks. She didn't advance.

Another candidate spent equal time on both — including a thoughtful take on how capability could be misused in education settings and what release safeguards she'd recommend. She did.

Come prepared with a mental model for this: Who benefits? Who could be harmed? What safeguards matter? How do you sequence a responsible rollout?

Step 4: The execution screen (metrics, but make it spicy)

Metrics. Root cause analysis. Debugging. The structure feels familiar but the content won't be.

The tradeoffs section is where people quietly fail this round. With AI products, tradeoffs aren't hypothetical case studies. They're live decisions affecting users, researchers, regulators, and the entire market.


Expect to get pushed. Hard. "Why this metric over that one?" "Why ship in 4 weeks instead of 8?" "You just cut this feature and what are you losing?"

Have opinions. Defend them. And be willing to update them when you get new information. Stubbornness is not conviction.

[See my answer to a recent Analytical question asked at Open AI.](#)

Step 5: The Loop (a.k.a. the main event)

Clear the screens and you enter the full loop. Here's every round, decoded.



The Parallel Team PM Round (30-45 min)

This is the most unique round in the process. It looks behavioral. It's actually a strategy conversation.

They want to know what you think. Not what frameworks you know. Not what considerations you'd weigh. What you actually, genuinely think the right strategy is

A candidate interviewing for a developer tools role got asked: "How do you think about expanding beyond early adopters into mainstream developers?"

What most candidates say: "Well, there are several considerations. You'd want to look at acquisition, activation, and retention across segments, and think about the developer experience..."

That's a consulting answer. It's technically correct and completely forgettable.

"I'd focus on reducing setup friction before scaling acquisition. Right now, trust is the bottleneck, not awareness. Developers have heard of the tool — they just don't believe it'll work in their stack without a 3-hour configuration headache. Fix that first, then pour fuel on growth."

See how that second answer has a spine? That's what they're looking for. A point of view you've actually thought through.

The Engineering Round

Standard collaboration and conflict questions, plus one curveball: they'll ask you to think about the technical roadmap.

What should OpenAI build? What are your assumptions about the current tech stack? Where's the leverage?

You don't need to write code. But "I'd defer to engineering on that" is a death sentence. Show that you engage with technical decisions. You have an informed perspective on what's buildable, what's hard, and where the technical bets should go

The Design Round

A case question specific to your area. If you're going for growth, expect something like: "What should a strong notifications strategy for ChatGPT look like?" If it's enterprise, maybe: "How would you think about customer segmentation?"

Plus behavioral questions about design collaboration. Pro tip: Don't tell a story where you were the hero who saved the project. Tell one that shows you understand what great PM-design partnership actually looks like. How you give feedback without being prescriptive. How you create space for design thinking instead of treating designers like pixel factories.

The Data Science Round

A case in your area plus behavioral. Growth candidates might get: "How would you measure success of a seat-bundling strategy for SMBs?" Enterprise candidates might hear: "How would you set metrics for a new collaborative chat feature?"

The behavioral piece is about your relationship with data. Not just "I defined north star metrics" (everyone says that). More like: How do you partner with DS to ask the right questions? What does a bad PM-DS relationship look like? Have you ever been wrong about what the data would show — and what did you do about it?

The Hiring Manager Round

Behavioral. This round is about you as a human. How you think. How you operate. What you care about.

By now, they've seen your product judgment and execution skills. This is about fit at a deeper level.

Be direct. Be honest about what you don't know. And have a real answer for what you do in your first 90 days — not a generic "listen and learn" answer, but a specific plan.

that shows you've already been thinking about the role like it's yours.

The Recruiter Debrief (the round nobody prepares for)

This is newer and genuinely unusual.

After the loop, you sit down with the recruiter and talk about how it went. What felt strong, what you'd do differently, any observations you have.


Most candidates assume the hard part is over. They wing it. Big mistake.

This is a self-awareness test. And it's a real evaluation round.

Think about it in advance: Which round was your best and why? Where did you struggle? If you had 10 more minutes in any round, what would you have added?

Answer honestly. Not with spin, not with false humility, not with "I think it all went great!" They can smell rehearsed self-deprecation from a mile away.

What actually separates the people who get offers



After working with dozens of candidates through this process, three things consistently separate offers from rejections:

Strong opinions, clearly defended. At every stage — produc

t, execution, strategy, HM — the winners know what they think. OpenAI isn't hiring professional fence-sitters. They're hiring people with conviction and the judgment back it up.

Obsessive preparation on the specific area. Generic prep is immediately obvious. The candidates who land offers know the product, the competitors, the open problems, they have a perspective on where things should go. They've done the work before the interview, not during it.

Honest self-reflection. This matters more than most people expect — especially in recruiter debrief. If you can evaluate your own performance clearly, without getting defensive or overly self-critical, that tells them you'll do the same thing on the job. And at a company with 30 PMs, everyone needs to be self-correcting.

The OpenAI PM bar is high. But it's not mysterious.

They want people who can think, who have a point of view, and who've done the actual work to prepare.

That's a higher bar than most companies. It's also a more honest one.

I coach PM candidates through OpenAI, Anthropic, Google, and other frontier AI interviews — from story and strategy to mock loops. If you're deep in prep and want a second set of eyes (or a first reality check), I'd love to help. Drop a comment or DM me.

Message Ridhima Khurana



18 Likes · 1 Restack

Discussion about this post

Comments

Restacks



Write a comment...

© 2026 Ridhima Khurana · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture